

MADHAV MERUVA

Applied AI & LLM Engineer | Backend & ML Systems

Hyderabad, India | +91 7670849767 | madhavmeruva690@gmail.com

GitHub: github.com/mad674 | **LinkedIn:** linkedin.com/in/madhavmeruva690 | **LeetCode:** leetcode.com/123mad

SUMMARY

Computer Science undergraduate specializing in **Applied AI, LLM systems, and backend engineering**, with hands-on experience building **semantic search** platforms, **retrieval-augmented** systems, and **end-to-end ML** pipelines. Experienced in developing production-ready AI applications using **LangChain, Vector Databases, PyTorch, TensorFlow, and FastAPI**. Built scalable AI-powered systems with **real-time inference, caching optimization, and modular API architectures**. Seeking Applied AI / ML Engineering internship opportunities.

EDUCATION

Keshav Memorial Institute of Technology (KMIT)

Bachelor of Technology, Computer Science Engineering | 2023 - 2027 | CGPA - 8.5

Sri Chaitanya Junior College

Intermediate (MPC) | 2021-2023 | GPA - 9.8

PROJECTS

AI Startup Idea Validator (LLM-Powered Evaluation Platform) | Aug 2025 – Oct 2025

Live: startupideavalidator.vercel.app | **GitHub:** github.com/mad674/startupideavalidator

- Built an LLM-powered startup evaluation platform integrating **semantic retrieval, vector search**, and AI-generated business insights.
- Implemented retrieval pipelines using **LangChain and PineconeDB** to improve contextual reasoning and relevance of generated responses.
- Improved inference performance by ~30% using **Redis caching, optimized API** orchestration, and scalable backend workflows.
- Designed **modular AI service architecture** supporting real-time interactions, retrieval-augmented generation, and concurrent user requests.

Technologies: React.js, Node.js (Express), FastAPI, MongoDB, PineconeDB, Redis, LangChain, OpenAI API, JWT

Financial Question Answering System (FinQA) | Feb 2025 - Jun 2025

Live: financeqptapp.netlify.app | **GitHub:** github.com/mad674/qa

- Built a **neuro-symbolic financial question-answering system** using a BERT retriever and LSTM-based program generator, and symbolic execution engine.
- Engineered **complete ML workflows** including preprocessing, model training, evaluation, and inference.
- Achieved **~80% accuracy** with interpretable program-step reasoning and domain-specific symbolic execution.

Technologies: Flutter, FastAPI, PyTorch, Firebase

Jewellery Design Pattern Generator (Pix2Pix GAN) | Aug 2024 - Dec 2024

Live: jewelrydesignpatternngen.netlify.app | **GitHub:** github.com/mad674/mlmodel

- Built a **Pix2Pix GAN-based system for sketch-to-image generation** with real-time inference capabilities.
- Developed a **real-time pipeline for generating design previews** from user-uploaded sketches.
- Improved output quality through **dataset preprocessing** and training optimization techniques.

Technologies: React.js, Node.js (Express), JWT, MongoDB, TensorFlow, Flask

SKILLS

- Languages:** Python, JavaScript, Java, C++
- AI/ML:** PyTorch, TensorFlow, NLP, Transformers, LangChain, Vector Databases
- Backend:** Node.js (Express), FastAPI, Flask, REST APIs
- Databases:** MongoDB, PineconeDB, Redis
- Cloud & DevOps:** Git, Postman, AWS, Docker
- Core Concepts:** Machine Learning, Deep Learning, DSA, OOP

ACHIEVEMENTS AND CERTIFICATES

- Solved **500+ DSA** problems across arrays, graphs, trees, and dynamic programming
- Secured a **top team award** at the KMIT Project Expo